

Architectural Test Task Group Call – Minutes

Thur, 22Jul2021 8am Pacific → Daylight ← Time

See slide 6 for agenda

Antitrust Policy Notice



RISC-V International meetings involve participation by industry competitors, and it is the intention of RISC-V International to conduct all its activities in accordance with applicable antitrust and competition laws. It is therefore extremely important that attendees adhere to meeting agendas, and be aware of, and not participate in, any activities that are prohibited under applicable US state, federal or foreign antitrust and competition laws.

Examples of types of actions that are prohibited at RISC-V International meetings and in connection with RISC-V International activities are described in the RISC-V International Regulations Article 7 available here: <https://riscv.org/regulations/>

If you have questions about these matters, please contact your company counsel.

RISC-V International Code of Conduct



RISC-V is a free and open ISA enabling a new era of processor innovation through open standard collaboration. Born in academia and research, RISC-V ISA delivers a new level of free, extensible software and hardware freedom on architecture, paving the way for the next 50 years of computing design and innovation.

We are a transparent, collaborative community where all are welcomed, and all members are encouraged to participate.

We as members, contributors, and leaders pledge to make participation in our community a harassment-free experience for everyone.

<https://riscv.org/risc-v-international-community-code-of-conduct/>

SIG Charter

The Architectural Compatibility Test SIG is an umbrella group that will provide guidance, strategy and oversight for the development of tests used to help find incompatibilities with the RISC-V Architecture as a step in the Architectural Compatibility self-certification process

The group will:

- Guide Development of:
 - Architectural tests for RISC-V implementations covering ratified and in-flight specifications for
 - Architectural versions, standard extensions, and implementation options.
 - Tools and infrastructure to help identify architectural incompatibilities in implementations
- Work with LSM and Chairs for resources to get the above work done.
- Mentor or arrange for mentoring for the resources to get the above work done

Administrative Pointers

- Chair – Allen Baum allen.baum@esperantotech.com Co-chair – Bill McSpadden bill.mcspadden@seagate.com
- SIG Email sig-arch-test@lists.riscv.org Notetakers: please send emails to allen.baum@esperantotech.com
- Meetings -Bi-monthly at 8am Pacific time on 2nd/4th Thursdays.
 - See https://docs.google.com/spreadsheets/d/1L15_gHI5b2ApkcHVtpZyl4s_A7sgSrNN zoom link

- Documents, calendar, roster, etc. in
 - <https://sites.google.com/a/riscv.org/riscv-staff/home/tech-groups-cal>
<https://drive.google.com/drive/folders/1DemKMAD3D0Ka1MeESRoVCJipSrwiUIEs>
 (lifecycle in "policies/supporting docs" folder, gaps in "planning" folder, arch-test specific in "information->content->arch-test")

Git repositories	← docs	riscv	→ tools
<ul style="list-style-type: none"> https://github.com/riscv/riscv-compliance/tree/master/doc/ https://riscv.readthedocs.io/en/latest/index.html https://riscv-isac.readthedocs.io/ https://riscv-ctg.readthedocs.io/ https://github.com/riscv/riscv-config/tree/master/docs https://github.com/remss-project/sail-riscv/tree/master/doc https://github.com/riscv-admin-docs/architecture-test/ 		tests riscv ISA coverage Test Gen. YAML, WARL config Sail formal model minutes, charter	https://github.com/riscv/riscv-arch-test/ https://gitlab.com/incoresemi/riscv/ https://github.com/riscv_isac https://github.com/riscv_ctg https://github.com/riscv/riscv-config/ https://github.com/remss-project/sail-riscv/

- JIRA: <https://jira.riscv.org/projects/CSC/issues/CSC-1?filter=allopenissues>
- Sail annotated ISA spec: in <https://github.com/remss-project/riscv-isa-manual/blob/sail/>
 - [README_SAIL](#) ← how to annotate annotated unpriv spec → [release/riscv-spec-sail-draft.pdf](#)
 - [release/riscv-spec-sail-draft.pdf](#) ← annotated source annotated priv spec → [release/riscv-privileged-sail-draft.pdf](#)
 - <https://us02web.zoom.us/rec/share/-XIYazzhIBbQoiZdarCfebdixjDWiVhf-LxnuVrliN4Bc30yf17ztKkKDU4Og54b.fArPPqnuR-NiXpQU> Tutorial

Access Passcode: tHAR#5\$V

Meeting Agenda

0. Looking for more admins, maintainers for riscv-arch-test git repo !!

I. Updates, Status, Progress:

1. Initial version of SAIL configurable using RISC-V-Config
2. Sail repo will be moving to riscv repo possibly as soon as next week
3. ACT policy delayed to modify errata policy
4. RISCOF migrated to github, Spike and Sail plugins, riscof branch is now in the main repo

II. Next steps and Ongoing maintenance

1. Discussion: other steps for Migration to Framework v.3.0 (riscov). (blocking items):

- a) Reorg of arch-test repo <done>
- b) Coverage reports: dev/vendor, archive in repo or not?
- c) (Sail/Spike model updates, pipecleaning, N people have run it, testing all the "fixed in riscov" issues
- d) Review Pipecleaner tests: What do we need to do to exercise capabilities for Priv Mode tests

2. Discussion: errata policy

3. Discussion: testing methodology for SIGs/TGs needing external stimulus/observability "ports".

4. Maintenance updates to V2 to enable future tests

- a) update RVTEST_SIGUPD to keep automatically adjust base/hidden offset when offset>2K,
- b) Enable use of Sail model results as the assertion value
- c) Convert assertions to be out-of-line
- d) add assertion macros for FP, DP, Vreg to arch_test.h and test_format spec
- e) add trap handlers for S, VS modes

5. Tests for non-deterministic result (see attached discussion in email)

- a) Provide a reference RTL test fixture (as opposed to SW functional model). See. JIRA CSC-6
- b) Define hooks for concurrency tests

6. Specific Arch-Test Policy/Process Gaps:

External Event ABI

- Why?
 - We want to be able to test events like: interrupts, concurrent reads & writes
 - These events would inject interrupts, modify memory at some future point
- What?
 - From a test perspective, these are model-specific macros that invoke vendor provided code
 - From an RTL perspective, this would look like a write to a specific MMIO “trick box” that RTL testbenches implements
- Possibilities
 - RVMODEL_ASSERT_INT(int, edgelevel, polarity, Trigger)
 - Int is a bitmask for simultaneous interrupts
 - Edgelevel, Polarity are masks for the type of signaling
 - Trigger is what initiates the event (e.g. #cycles from the write, debug trigger event, instret offset)
 - RVMODEL_MEMWRT(address, data, event)
- Questions
 - Should there be a pseudo-randomized Trigger?
 - What events should be standardized?
 - Do we need deassertion macros, Read macros? (we actually have specific interrupt deassertion macros now).
 - Are there another type of Event we should consider?

Discussion

Sail Status: Initial version of YAML configurable Sail created (YAML->python->Sail)
Sail repo now has a maintainer

RISCOF migration status

RISCOF repo migrated to github and updated : <https://github.com/riscv/riscof>.

- new command [arch-tests](#) added to RISCOF; it clones latest (or specific) riscv-arch-test repo version to avoid need to clone it manually.

- RISCOF [docs](#) for have been updated & other clean up was done; changelog here:

<https://github.com/riscv/riscof/blob/master/CHANGELOG.md#1210---2021-07-21>

Spike DUT plugin PR created : <https://github.com/riscv/riscv-isa-sim/pull/748>

Sail Ref plugin PR created : <https://github.com/rem-s-project/sail-riscv/pull/97>

riscv-dev branch on riscv-arch-test repo created here:

<https://github.com/riscv/riscv-arch-test/tree/riscv-dev>

This removes the following

- top-level Makefile and Makefile.include

- old riscv-test-env which were meant for migration
- ovpsim directory

- All riscv-targets which were compatible only with v2.
- migration guide from doc/ directory

- riscv-test-stats directory retained until we figure out what we want to do with them

- Porting guide in doc/README.adoc now points to RISCOF's guide.

(This only has the vocabulary which now is relevant to riscv porting)

- Suggest that the new riscv-dev branch merge doc/README.adoc contents to the top-level README.md. This would include copying the following sections:

+Intent +vocabulary +directory structure

Coverage reports

Chair: Coverage reports are artifacts. How to handle? Users put test results into github directory for self-certification, but will also get coverage. User may not implement all CSR functionality, so coverage won't match test developer coverage.

Inspire: Are you saying that some things aren't run?

Chair: No. E.g. some optional registers or CSR bit are not implemented;

Extensions must be fully covered. Reports from test run & passed/fail and coverage

– which won't see CSR bits that aren't implemented

– so coverage will look incomplete, but really aren't from the models perspective

There is a vendor results area that a user can put their results. For self-attestation.

InCore: Coverage report is derived from SAIL execution, so novendor-specific info

CTO: question: for branding, what do we say about coverage reports?

Chair: Today, nothing. It isn't considered.

CTO: Do tests today fail if someone has not implemented an instruction?

Chair: If required for an extension that is claimed to be supported, yes.

CTO: assertion: coverage report is not needed for branding.

for that, we require only: 1) pass/fail report 2) errata

Chair: Agreed; then do we even put vendor coverage report into repo
(coverage reports that WE generate, WILL go into the repo)

CTO: perhaps leave it to the vendor? May be useful for the vendor comparison

Chair: They can do that without putting into the test reports repo..

Chair: **Decision:** vendors won't add their coverage reports to test-report repo

Errata policy

CTO: This needs to be dealt with situation Krste talked about at tech-chairs yesterday

Categories: test problems, spec problem, config problem, implementation problem

if vendor fails ACT, then custom. can't brand without a waiver

if vendor passes ACT, but they still have a failure, then errata.

(There may never be an ACT to cover the particular failure)

if vendor passes ACT, but new test causes failure: can claim compatibility for that version w/ errata

if they spin the chip and don't fix it, then they are a "custom SOC".

There needs to be some slop in identifying status

InCore: but we're running on RTL. not silicon.

CTO: let's move to offline < discussion of timeline. 6 months slop proposed. self-governed >

Public peer process will keep people honest

Chair: policy needs to be reviewed carefully by vendors: see

https://docs.google.com/document/d/1ZKSLQ5HPT3E_CqTVOQIBPs7qJ9v1mpnMDXHhtOQJFcU/

CTO: certification /attestation pushed to marketing. We only give basic input on compatibility

External Event ABI: see slide 7

Chair: Some things can't be tested because we need external stimulus. ex: external interrupts.

Cochair: - need open source model, and proceed in two steps

1st : all interrupts were serviced, 2nd - did they happen in the right order

CTO: restrict test handlers not to modify tests. < need to write down limitations of interrupt >

Chair: ex: ASSERT_INT(intID, type, offset, trigger)

Are pseudo random offsets even useful? We would need to keep same seed between ref and

every model, and maintain ordering. It may as well be a constant; But what is offset unit: cycles?

insts retired? What is trigger based on? Issue addr? Instret (would need to be >=comparison)

InCore: will be model specific

Chair: We need to pick something to get as close as possible for ordering events.

Decisions & Action Items

Decisions

- Riscov will be added as an in an arch-test branch ahead of the switch over (all existing tests will be backwards compatible)
- separately, a branch of the arch-test repo will be created with all the directory changes needed for the switchover to riscov.
- Riscov plugins will be moved to model repos & sample plugins will be created for Sail and Spike
- all READMEs will be updated to indicate when the switchover will happen, and point to the riscov branch and encourage/warn users to start using it
- In (very) roughly 3 months, we will merge in branch.
- Coverage reports generated as a byproduct of tests merged into the test suite will be located in the riscv-arch-test/riscv-tests-suite github directory folder; vendor coverage reports will not, and are only for the benefit of the vendor.

Outstanding Action Items

- create framework riscv repo <done>
- Add branch for reconfigured riscov arch-test repo & migrate riscov to new framework repo <done>
- Fix uses of RVTEST_ISA macro in various tests (formatting incompatibility with riscov and update spec <Neel>
- Contact SW HC & DOC SIG to determine an inline comment->doc tool flow, and determine if docs (as opposed to ISA specs) must be .adoc, or could be .pdf or .html <Allen, Jeff-in progress>
- Settle where plugins will go, and the migration process <done>
- Update all READMEs to point to branch <Neel, Pawan?>
- Set agenda item to decide how to handle coverage reports <done>
- Update standard trap handler code for added priv levels, custom exception handler registration, <Allen>

Pull/Issue Status

Issue#	Date	submitter	title	status	comments
#4	03-Jul-2018	Kasanovic	Section 2.3 Target Environment	Fixed in riscov	Will be closed in V3
#22	24-Nov-18	brouhaha	I-MISALIGN_LDST-01 assumes misaligned data access will trap	^	HW misalign support not configurable now
#40	4-Feb-19	debs-sifive	Usage of tohost/fromhost should be removed		Not RV32EC or RV32EM
#142	17-Nov-20	subhajit26	Not able to run compliance test for rv32E device and RV32E ISA	RV32E only	HW misalign support not configurable
#146-9	01-Dec-20	Imperas	Test I EBREAK,ECALL, MISALIGN_JMP/LDST, OpenHW	v	Close on 31Jul unless objections
#193	18-Jun-21	odarcy1	CHECK_XLEN macro will silently pass	N/A	-will close in 3wks; all uses deprecated
#107	22-Apr-20	jeremybennett	Clang/LLVM doesn't support all CSRs used in compliance test suite	under discussion	Who can review this?
#115	06-jun-20	adchd	How to support on-board execution?	under discussion	Approved, needs merge
pull#129	31-jul-20	nmeum	sail-riscv-ocaml: Disable RVC extension on all devices not using it	In process	
pull#184	15-apr-21	dansmathers	Updating http reference for constr	In process	
#119	17-jun-20	allenjbaum	Missing RV32i/RV64i test: Fence	Test has been written	Close when RFQ test is merged
#188	26-Apr-21	neelgala	Updates required in K_unratified tests to be compatible with current RISCOF		
#189	26-Apr-21	neelgala	Proposal to enhance the RVTEST_ISA macro		
#190	26-Apr-21	neelgala	The 16-byte signature boundary issue		

JIRA Status

Issue#	Date	submitter	title	status	comments
IT-1	27Aug/20	Allen Baum	Need to modify the description of compliance in https://riscv.org/technical/specifications/	done	
IT-4	01/Sep/20	Allen Baum	Add Jira link to TG home pages	done	
CSC-1	20/Aug/20	Ken Dockser	Come up with names for the tests suites that we are creating		1 st step done
CSC-2	20/Aug/20	Ken Dockser	Produce concise text to explain the Architecture Tests intent and Limits	done	Written, needs pull req
CSC-3	20/Aug/20	Ken Dockser	Come up with an internal goal for what we wish to accomplish with the Architectural Tests		Not written
CSC-4	20/Aug/20	Ken Dockser	Develop a roadmap for all the different categories of test suites that will need to be created		Not written
CSC-5	20/Aug/20	Ken Dockser	Develop a roadmap for releases of single-instruction Architecture Tests		Not written
CSC-6	20/Aug/20	Ken Dockser	Develop a reference RTL test fixture that can stimulate and check the CPU under test		Needs more discussion

BACKUP

Test Acceptance Criteria

Tests merged into the ACT test_suite repo must :

- conform to the current format spec (macros, labels, directory structure)
 - including framework-readable configurations - i.e. which ISA extension it will be tested with (using Test_Case macro parameter equations) for each test case
- use only files that are part of the defined support files in the repository, including standard trap handlers
 - TBD: how to install test specific (not model specific) handlers
- Be able to be loaded, initialized, run, signal completion, and have signature results extracted from memory by a/the framework
- run using the SAIL model and not fail any tests
- generate signature values either
 - directly from an instruction result (that can be saved & compared with DUT/sim)
 - by comparing an instruction result with a configuration-independent value range embedded in the test code (e.g. saving above, below, within)
 - by comparing an instruction result with a configuration-independent list of values (e.g saving matches or mismatched)
 - (it can be useful to also return a histogram of value indices that matched)
- Store each signature value into a unique memory location in a signature region that is
 - delimited by standard macros embedded in the test which can be communicated to the test framework
 - pre-initialized to values that are guaranteed not to be produced by a test
- have defined coverage goals in a machine readable form that can be mechanically verified
- improve coverage (compared to existing tests) as measured and reported by a coverage tool (e.g. ISAC)
- use only standard instructions (and fixed size per architecture macros, e.g. LI, LA are allowed)
- be commented in test_case header (ideally listing coverpoint covered)

Tests that are otherwise accepted, but depend on tools or simulators that have not be upstreamed must be put into a <Ext-Name_unratified>/ directory instead of <Ext-Name>/

Framework Requirements – first cut

The framework must:

- Use the TestFormat spec and macros described therein
 - (which must work - including assertions)
- Choose test cases according to equations that reference the YAML configuration
- Define macro variables to be used inside tests based on the YAML configuration
- Include the compliance trap handler(s), & handle its (separate) signature area(s)
- Load, initialize, and run selected tests between two selected models, extract the signatures, compare results, and write out a report file
- Exist in a riscv github repo, with a more than one maintainer.
- Be easy to get running, e.g.:
 - run under a variety of OSes with the minimum number of distro specific tools.
 - Not require sudo privileges
- Have the ability to measure and report coverage for test generation
 - Coverage specification is a separate file
 - Could be a separate app

Non-determinism in Architectural Tests

The RV architecture defines optional and model/ μ arch defined behavior.

This implication: there are tests that have multiple correct answers. E.g.:

- Misaligned accesses: can be handled in HW, by "invisible" traps w/ either misaligned or illegal access causes, and do it differently for the same op accessing the same address at different times (e.g. if the 2nd half was in the TLB or not)
- Unordered Vector Reduce ops: (different results depending on ordering & cancellation)
- Tests involving concurrency will have different results depending on microarchitectural state, speculation, or timing between concurrent threads (e.g. modifying page table entry without fencing)

From the point of view of ACTs, there are 2 (& sometimes more) legal answers. The golden model only generates one. Possible mechanisms to test include:

- Modify (if necessary) & configure reference model to generate each legal result, run it with each config, & accept either result from the DUT (e.g. misalign or un-fenced PTE modification)
- Provide specific handlers for optional traps
- Use self-testing tests(compare with list or range of allowed outcomes from litmus tests)
- Avoid tests that can generate non-deterministic results
- Ultimately: develop new frameworks that can handle concurrency along with reference models that can generate all legal outcomes
- It is the responsibility of the TG that develops an extension to develop the strategy for testing features and extensions that can have nondeterministic results

TGs under the SIG

- IF you're creating work product, you should be a TG
- If changing requirements, plans ABIs, etc
 - Test plan==SOW
- The Architectural Compatibility Test Task Group will define and maintain specifications for
 - test formats
 - test-benches and frameworks needed for
 - privilege testing privilege testing,
 - Concurrency/ Memory model testing
 - Asynchronous event testing (interrupts)
 - Nondeterministic tests
 - ISA test coverage goals
 - test tools (e.g. coverage, generators)
- The Architectural Compatibility Test Task Group will maintain the appropriate GitHub:
 - tests for the individual ISA extensions
 - issues related to the tests
 - the operation and issues related to the framework
- The Architectural Compatibility Test Task Group will
 - work with the different privilege and un-privilege ISA extension Task Groups
 - to help them write test plans/specs for the ISA tests
 - to help them work with the sub-contractors (IITMadras, RIOS, CAS, etc) to deliver the tests
 - assess quality of delivered tests and be maintainer for the test GitHub

Meeting Conventions



- We don't solve problems or detailed topics in most meetings unless specified in the agenda because we don't often have enough time to do so and it is more efficient to do so offline and/or in email. We identify items and send folks off to do the work and come back with solutions or proposals.
- If some policy, org, extension, etc. can be doing things in a better way, help us make it better. Do not change or not abide by the item unilaterally. Instead let's work together to make it better.
- Please conduct meetings that accommodates the virtual and broad geographical nature of our teams. This includes meeting times, repeating questions before you answer, at appropriate times polling attendees, guide people to interact in a way that has attendees taking turns speaking, ...