Architectural Test SIG Call –Minutes

Thur, 14Oct2021 8am Pacific → Daylight ← Time

See slide 6 for agenda

RISC-V attendance

Only RISC-V Members May Attend

- Non-members are asked to please leave.
- Members share IP protection by virtue of their common membership agreement. Non-members being present jeopardizes that protection
- It is easy to become a member. Check out riscv.org/membership
- If you need work done between non-members or or other orgs and RISC-V, please use a joint working group (JWG).
  - used to allow non-members in SIGs but the SIGs purpose has changed.
- Please put your name and company (in parens after your name) as your zoom name. If you are an individual member just use the word “individual” instead of company name.
- Non-member guests may present to the group but should only stay for the presentation. Guests should leave for any follow on discussions.

RISC-V®

Antitrust Policy Notice

RISC-V International meetings involve participation by industry competitors, and it is the intention of RISC-V International to conduct all its activities in accordance with applicable antitrust and competition laws. It is therefore extremely important that attendees adhere to meeting agendas, and be aware of, and not participate in, any activities that are prohibited under applicable US state, federal or foreign antitrust and competition laws.

Examples of types of actions that are prohibited at RISC-V International meetings and in connection with RISC-V International activities are described in the RISC-V International Regulations Article 7 available here: https://riscv.org/regulations/

If you have questions about these matters, please contact your company counsel.

RISC-V®

Collaborative & Welcoming Community

RISC-V is a free and open ISA enabling a new era of processor innovation through open standard collaboration. Born in academia and research, RISC-V ISA delivers a new level of free, extensible software and hardware freedom on architecture, paving the way for the next 50 years of computing design and innovation.

We are a transparent, collaborative community where all are welcomed, and all members are encouraged to participate. We are a continuous improvement organization. If you see something that can be improved, please tell us. help@riscv.org

We as members, contributors, and leaders pledge to make participation in our community a harassment-free experience for everyone.

https://riscv.org/community/community-code-of-conduct/

RISC-V®

SIG Charter

The Architectural Compatibility Test SIG is an umbrella group that will provide guidance, strategy and oversight for the development of tests used to help find incompatibilities with the RISC-V Architecture as a step in the Architectural Compatibility self-certification process

The group will:

- Guide Development of:
  - Architectural tests for RISC-V implementations covering ratified and in-flight specifications for
    - Architectural versions, standard extensions, and implementation options.
  - Tools and infrastructure to help identify architectural incompatibilities in implementations
- Work with LSM and Chairs for resources to get the above work done.
- Mentor or arrange for mentoring for the resources to get the above work done

Adminstrative Pointers

- Chair – Allen Baum allen.baum@esperantotech.com Co-chair – Bill McSpadden bill.mcspadden@seagate.com
- SIG Email sig-arch-test@lists.riscv.org. Notetakers: please send emails to allen.baum@esperantotech.com
- Meetings -Bi-monthly at 8am Pacific time on 2nd/4th Thursdays.
  - See https://docs.google.com/spreadsheets/d/1L15_gHl5b2ApkcHVtpZyl4s_A7sgSrNN zoom link
- Documents, calendar, roster, etc. in
  - https://sites.google.com/a/riscv.org/risc-v-staff/home/tech-groups-cal
  - https://drive.google.com/drive/folders/1DemKMAD3D0Ka1MeESRoVCJipSrwiUlEs

    lifecycle in "policies/supporting docs" folder, gaps in "planning" folder, arch-test specific in "information->content->arch-test")
- Git repositories ←docs riscv → tools

←docs	riscv	→ tools
https://github.com/riscv/riscv-compliance/tree/master/doc/	tests	https://github.com/riscv/riscv-arch-test/
https://riscof.readthedocs.io/en/stable/	riscof	https://github.com/riscv/riscof
https://riscv-isac.readthedocs.io/	ISA coverage	https://github.com/riscv_isac
https://riscv-ctg.readthedocs.io/	Test Gen.	https://github.com/riscv_ctg
https://github.com/riscv/riscv-config/tree/master/docs	YAML, WARL config	https://github.com/riscv/riscv-config/
https://github.com/riscv/sail-riscv/tree/master/doc	Sail formal model	https://github.com/riscv/sail-riscv/
https://github.com/riscv-admin/architecture-test	minutes, charter	

- JIRA: https://jira.riscv.org/projects/CSC/issues/CSC-1?filter=allopenissues
- Sail annotated ISA spec: in https://github.com/rems-project/riscv-isa-manual/blob/sail/

README.SAIL ←how to annotate	annotated unpriv spec→ release/riscv-spec-sail-draft.pdf
release/riscv-spec-sail-draft.pdf ← annotated source	annotated priv spec→ release/riscv-privileged-sail-draft.pdf
https://us02web.zoom.us/rec/share/-XIYazzhIBbQoiZdarCfebdjxjDWiVhf-LxnuVrliN4Bc30yf17ztKkKDU4Og54b.fArPPqnuR-NiXpQU Tutorial Passcode: tHAR#5$V	

Meeting Agenda

0. **Looking for more admins, maintainers for riscv-arch-test git repo !!**

I. Updates, Status, Progress:

I. RV32F merged, RV64D pull request made, RV32D now supported in Sail, RV64F coming along

II. BitManip tests almost ready, Sail support pull request filed

II. Next steps and Ongoing maintenance

1. Discussion: other steps for Migration to Framework v.3.0 (riscof). (blocking items):
   a) Marc (Inspire) will give a brief update on integrating the riscof framework their TB/verification environment.
   b) (Sail/Spike model updates, pipecleaning, N people have run it, testing all the "fixed in riscof" issues
   c) Review Pipecleaner tests: What do we need to do to exercise capabilities for Priv Mode tests
2. Dynamic Test Generation
   1. Related: how should we deal with 1GB test directories (FP
3. Discussion: starting a TG to precisely define the ABI for asynch event generation and a reference C-model to be used for Sail
4. Config YAML GUI interface demo

III. Future Agenda items

1. Maintenance updates to V2 to enable future tests
   a) update RVTEST_SIGUPD to keep automatically adjust base/hidden offset when offset>2K, - ongoing
   b) Enable use of Sail model results as the assertion value -
   c) Convert assertions to be out-of-line
   d) add assertion macros for FP, DP, Vreg to arch_test.h and test_format spec
   e) add trap handlers for S, VS modes
2. Tests for non-deterministic result (see attached discussion in email)
   a) Provide a reference RTL test fixture (as opposed to SW functional model). See. JIRA CSC-6
   b) Define hooks for concurrency tests

Discussion

Status:
Inspire: have an example RISCOF installation located at
https://github.com/InspireSemi/riscof_install_example_DUT/
(see next slide). It does expect that the gnu toolchain is installed.
We are running arch-tests on RV64IMCF under on RHEL8
Chair: which needs sudo privs to install, and which can be installed in home directory?
Inspire: Dev Tools and Z3 package need sudo
I would not recommend that anyone try and run RISCOF on RH7 or CentOS7.
Chair: Is this because of SAIL?
Inspire: it's because of a lot of stuff; cmake is one (fixed w/ Sail docker image)
Chair : Do you need sudo do you install docker?
Or can you use rootless docker engine?
Inspire: That requires downloading a lot of source, compiling them all and all their dependencies, which themselves need sudo privs; still an admin issue.
<walking through local_install_sail_rhel8.sh>

Cochair: how long to run it in SAIL and get signatures?
Incore: about a minute
CTO: smoke tests should be separated out. for CI needs to be fast but runs on sims.
POC doesn't need to be open source, Can run tests in parallel: for each ELF file.
CTO: CI can't run for hours
what about creating a smoke test suite to run?

Chair: running individual tests means creating and editing test reports; we would prefer if that wasn't even possible to prevent altering results
CTO: reports are to be self-attested; we don't need that level of control.

PQShield: I heard about YAML and config in this discussion. Where does this fit into the RV config?
< discussion about config-struct >
Seagate: We need to ensure no spec divergence between what we are told to test (risc-config YAML) and what the core reports it is implementing (config-struct ASN.1)
CTO: I've been driving to get convergence: ASN1 and config YAML.

IA: config structure needs translators to/from ACT YAML.
IA**:** Demo a YAML GUI interface at the next meeting which eases burden on implementers.

Example riscof repo

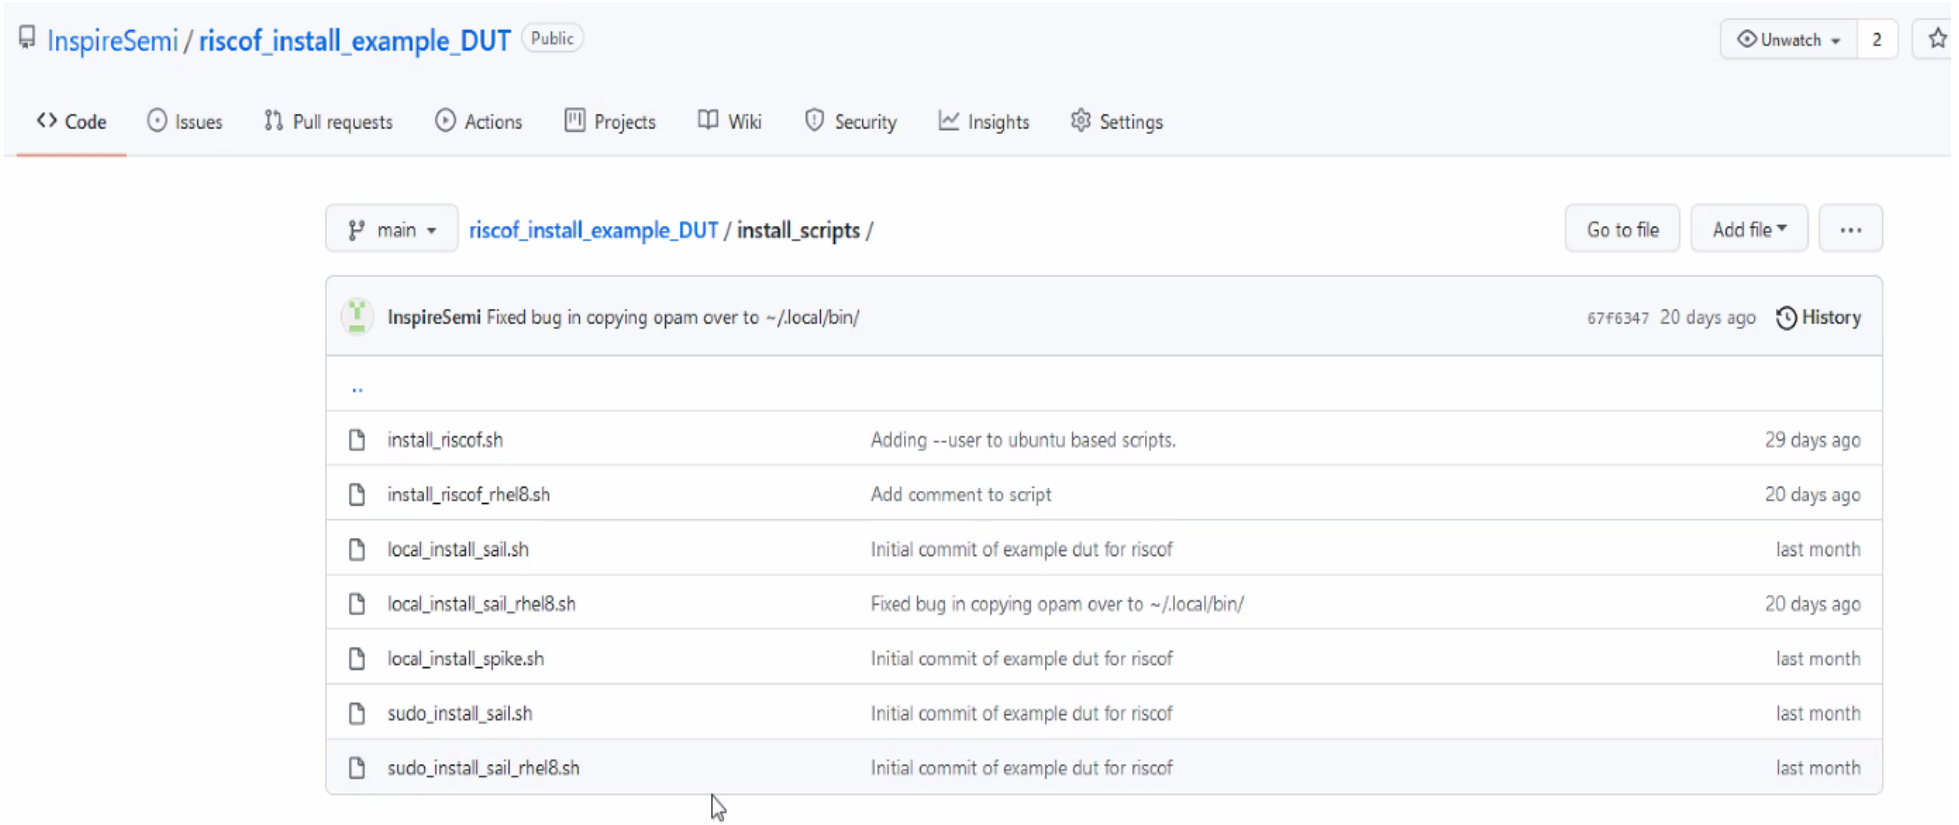

Decisions & Action Items

Decisions (from last meeting)

- Revised Decision : Sail interface to Asynch will be an MMIO write
  - Address is implementation dependent.
    Easier to implement in RTL
- Decision : Sail will trigger event generation when instret > delay+(instret@Store of) Delay

Outstanding Action Items

- config structure needs translators to/from ACT YAML <added to Dev partners spreadsheet>
- Demo a YAML GUI interface at the next meeting which eases burden on implementers <nezt meeting>
- Develop a C-model of an asynch event generator and a Sail+Spike interface to it
- Update all READMEs to point to branch <Neel, Pawan?>
- Update standard trap handler code for added priv levels, custom exception handler registration, <Allen, under review>
- Contact SW HC & DOC SIG to determine an inline comment->doc tool flow, and determine if docs (as opposed to ISA specs) must be .adoc, or could be .pdf or .hmtl <Allen, Jeff-in progress>
- Marc's example plugin to be added to riscof repo <Marc,Neel> (with updated documentation)

BACKUP

External Asynch Event Support

- Why?
  - We want to be able to test events like: interrupts, concurrent reads & writes
  - These events would inject interrupts (wired and msg signaled), modify memory
- What?
  - From at test perspective, these are model-specific macros that invoke vendor provided code
  - From an RTL perspective, this would look like a write to a specific MMIO “trick box” that RTL testbenches implement
    - We will provide one as a Sail external function in C, and Spike
  - The ACT SIG (or a TG under it) will develop an implementation template as a C model
- Strawman C interface
  - RVMODEL_EVENT(instret, delta, eventtype_num)
  - RVMODEL_MEMWRT(instret, delta, size, address data) ←note that this could be a Message Signalled interrupt
  - RVMODEL_MEMRD(instret, delta, size, address)
    - Assembly language implementation are 1-3 successive writes; events are scheduled on the last one
    - Vendor code is called every cycle, and asserts events to the model when curr instret>=orig_instret+delta (instret is global variable, sampled on final MMIO write)
    - Possible schedule multiple events with appropriate deltas, even enabling simultaneous events + rd/write} (but only single rd/wr per cycle)
  - RVMODEL_EVT_STATUS() – returns values read by RVMODEL_MEMRD()
- Strawman ACT Macro Interface
  - RVMODEL_EVENT(typ_evtnum_ delta_srcreg)
    - Assembly language would store 1 word at a platform specific address/CSR: edgelvl/polaritytype/evt_num encoded in high half, delta in lower)
  - RVMODEL_MEMWRT(addr_srcreg, data_srcreg, sz_delta_srcreg) ←note that this could be a Message Signaled interrupt
    - Assembly language would store 3 words at platform specific mmio addresses/CSRs: address, data, sz+delta
  - RVMODEL_MEMRD(addr_srcreg,typ_ sz_delta_srcreg)
    - Assembly language would store 3 words at platform specific address or into 3 CSRs: address, data, type+sz+delta (type is ownership/temporal, etc)

Pull/Issue Status

Issue#	Date	submitter	title	status	comments
#4	03-Jul-2018	Kasanovic	Section 2.3 Target Environment	Fixed in riscof	Will be closed in V3
#22	24-Nov-18	brouhaha	I-MISALIGN_LDST-01 assumes misaligned data access will trap	^	HW misalign support not configurable
#40	4-Feb-19	debs-sifive	Usage of tohost/fromhost should be removed	|	now
#146-9	01-Dec-20	Imperas	Test I EBREAK,ECALL, MISALIGN_JMP/LDST, OpenHW	v	HW misalign support not configurable
#115	06-jun-20	adchd	How to support on-board execution?	under discussion	
pull#129	31-jul-20	nmeum	sail-riscv-ocaml: Disable RVC extension on all devices not using it	In process	Who can review this?
pull#184	15-apr-21	dansmathers	Updating http reference for constr	In process	Approved, needs merge
pull#199	01-Aug-21	bilalsakhawat	Fix for issue #142 , Adds RV32EC, EM tests		Wait for RV32E spec? rename unratified
#119	17-jun-20	allenjbaum	Missing RV32i/RV64i test: Fence	Test has been written	Close when RFQ test is merged
#189	26-Apr-21	neelgala	Proposal to enhance the RVTEST_ISA macro		
#190	26-Apr-21	neelgala	The 16-byte signature boundary issue		
Pull#201	17-Aug-21	Liweiwei90	Update K-ext tests		Updates for spec changes, improved Sbox coverage
#203	24-Aug-21	Allenjbaum	Fence test has poor coverage		Specifically: test fm bits are ignored

JIRA Status

Issue#	Date	submitter	title	status	comments
IT-1	27Aug/20	Allen Baum	Need to modify the description of compliance in https://riscv.org/technical/specifications/	done	
IT-4	01/Sep/20	Allen Baum	Add Jira link to TG home pages	done	
CSC-1	20/Aug/20	Ken Dockser	Come up with names for the tests suites that we are creating		1st step done
CSC-2	20/Aug/20	Ken Dockser	Produce concise text to explain the Architecture Tests intent and Limits	done	Will become ACT policy
CSC-3	20/Aug/20	Ken Dockser	Come up with an internal goal for what we wish to accomplish with the Architectural Tests		Will become ACT policy
CSC-4	20/Aug/20	Ken Dockser	Develop a roadmap for all the different categories of test suites that will need to be created		Not written
CSC-5	20/Aug/20	Ken Dockser	Develop a roadmap for releases of single-instruction Architecture Tests		Not written
CSC-6	20/Aug/20	Ken Dockser	Develop a reference RTL test fixture that can stimulate and check the CPU under test		Needs more discussion

Non-determinism in Architectural Tests

The RV architecture defines optional and model/µarch defined behavior. This implication: there are tests that have multiple correct answers. E.g.:

- Misaligned accesses: can be handled in HW, by "invisible" traps w/ either misaligned or illegal access causes, and do it differently for the same op accessing the same address at different times (e.g. if the 2nd half was in the TLB or not)
- Unordered Vector Reduce ops: (different results depending on ordering & cancellation)
- Tests involving concurrency will have different results depending on microarchitectural state, speculation, or timing between concurrent threads (e.g. modifying page table entry without fencing)

From the point of view of ACTs, there are 2 (& sometimes more) legal answers. The golden model only generates one. Possible mechanisms to test include:

- Modify (if necessary) & configure reference model to generate each legal result, run it with each config, & accept either result from the DUT (e.g. misalign or un-fenced PTE modification)
- Provide specific handlers for optional traps
- Use self-testing tests(compare with list or range of allowed outcomes from litmus tests)
- Avoid tests that can generate non-deterministic results
- Ultimately: develop new frameworks that can handle concurrency along with reference models that can generate all legal outcomes
- It is the responsibility of the TG that develops an extension to develop the strategy for testing features and extensions that can have nondeterministic results

Framework Requirements

The framework must:

- Use the TestFormat spec and macros described therein
  - (which must work - including assertions)
- Choose test cases according to equations that reference the YAML configuration
- Define macro variables to be used inside tests based on the YAML configuration
- Include the compliance trap handler(s), & handle its (separate) signature area(s)
- Load, initialize, and run selected tests between two selected models, extract the signatures, compare results, and write out a report file
- Exist in a riscv github repo, with a more than one maintainer.
- Be easy to get running, e.g.:
  - run under a variety of OSes with the minimum number of distro specific tools.
  - Not require sudo privileges
- Have the ability to measure and report coverage for test generation
  - Coverage specification is a separate file
  - Could be a separate app

Test Acceptance Criteria

Tests merged into the ACT test_suite repo must :

- conform to the current format spec (macros, labels, directory structure)
  - including framework-readable configurations - i.e. which ISA extension it will be tested with (using Test_Case macro parameter equations) for each test case
- use only files that are part of the defined support files in the repository, including standard trap handlers
  - TBD: how to install test specific (not model specific) handlers
- Be able to be loaded, initialized, run, signal completion, and have signature results extracted from memory by a/the framework
- run using the SAIL model and not fail any tests
- generate signature values either
  - directly from an instruction result (that can be saved & compared with DUT/sim)
  - by comparing an instruction result with a configuration-independent value range embedded in the test code (e.g. saving above, below, within)
  - by comparing an instruction result with a configuration-independent list of values (e.g saving matches or mismatched)
    - (it can be useful to also return a histogram of value indices that matched)
- Store each signature value into a unique memory location in a signature region that is
  - delimited by standard macros embedded in the test which can be communicated to the test framework
  - pre-initialized to values that are guaranteed not to be produced by a test
- have defined coverage goals in a machine readable form that can be mechanically verified
- improve coverage (compared to existing tests) as measured and reported by a coverage tool (e.g. ISAC)
- use only standard instructions (and fixed size per architecture macros, e.g. LI, LA are allowed)
- be commented in test_case header (ideally listing coverpoint covered)

Tests that are otherwise accepted, but depend on tools or simulators that have not be upstreamed must be put into a <Ext-Name_unratified>/ directory instead of <Ext-Name>/